CS 7641 Machine Learning
Assignment 1

Philip Bale
pbale3

Due February 4th, 2018 11:59pm

Classification Problems

1) US Permanent Visa Applications

Overview

The first classification problem revolves around classifying whether or not a person's US permanent visa application will be accepted or denied based on the parameters of their application. Among the features used in the classifcation are:

- Job features: Industry code, job class, wage rate, wage type

- Geographic features: Country of citizen and employer location

The classes observed for this dataset are simply 'approved' and 'denied'. The dataset contains 374365 total samples.

Why is the dataset interesting?

This dataset is interesting due to its potential to aid in the visa application process from a cost and time savings potential. It could also enable confidence in those interested in applying for a US permanent visa but doubting their chances of acceptance. At the end of the day, the goal is it to try to determine the application result before time, money, and other resources are spent. As someone who has worked with a large number of first-generation visa holders and immigrants, I am extremely interested in building tools to help others to achieve the same.

From a machine learning perspective, the dataset is incredibly interesting due to its wide variety of features and the variety of values those features can take. An immense number of job types, wage rates, and citizenships alone create an extremely diverse dataset. Additionally, the number of samples available provide a comprehensive picture of historical data, lending towards greater cofidence in training and testing rates.

2) Home Sale Price Predictions

Overview

The second classification problem revolves around classifying a home's price bracket based upon the various characteristics of the home. Among the features used in the classification are :

- Subjective measurements: Exterior condition, house style, overall quality rating, and overall condition

- Objective measurements: Type of dwelling, building type, lot size, neighborhood, year built, and year sold

After an initial review of the dataset, the classes were defined as pricing brackets divided into 100k groups. I.e: 0-100k, 100k-200k, 200k-300k, etc. The dataset contains 1451 samples. An additional dataset containing another 1400 testing samples exists but was not used as it contains unclassified sale prices. It will, however, prove useful for unsupervised learning.

Why is the dataset interesting?

This dataset is interesting for two primary reasons: real-world applicability and participating in a Kaggle challenge. First, modeling home prices is both a difficult and lucrative task. If one can succesfully model home sale prices on large sets of data, he/she can make large amounts of money investing in real estate when he/she detects outliers in listed price vs. what it is expected to sell for. This applies to flipping, investing, and remodeling. Second, the dataset is part of an ongoing Kaggle competition that does not have a winning solution yet. By taking part of the competition, the dataset presents the

opportunity to work towards a winning solution and advance ones algorithms over time.

Houses can have a very large amount of features–with a large amount of variety in the individual features. Similarly, housing is prone to personal taste and frequent need for upgrades/modernization. In such, I believe price estimation is an excellent problem, full of depth and complexity, that is suitable for a machine learning approach.

General Data Processing

The datasets I used were both relatively clean to begin with. One small problem, however, was that a lot of my features on both datasets were text-based. To transform the features into numeric values suitable for the machine learning algorithms, I used a label encoder built into ScikitLearn.

I also did a small amount of preprocessing of the data to make it more suitable for classification. I dropped all unnecessary columns to help speed up with data processing in general–which proved immensely helpful when dealing with the more computationall intensive algorithms. In the case of home prices, I precalculated the brackets based on the 'sale price' data label. In the case of visa applications, I pregrouped the case outcome so that results such as 'certified' and 'certified withdrawn' are both concerned as 'approved' conditions whereas 'denied', 'invalidated', and 'rejected' all resolve to 'denied'.

1: Decision Trees

A decision tree classifier was the first algorithm applied to the datasets. Various values of max_depth were tested as a means of pruning unnecessary leaves. Similarly, a grid search was used to test whether a 'gini' or 'entropy' criterion was more effective.

US Permanent Visa Data

Depth	Criterion	Tree Size	Train %	Train Time	Test %	Test Time
1	gini	3	0.7657	0.1212	0.7680	0.0010
3	gini	15	0.6775	0.1417	0.6749	0.0009
6	entropy	105	0.7121	0.2605	0.6924	0.0011
10	gini	889	0.7584	0.3192	0.7072	0.0011
15	gini	3013	0.8628	0.4199	0.7582	0.0014
20	gini	4751	0.9299	0.4517	0.7917	0.0010
25	gini	5349	0.9545	0.5284	0.8032	0.0010
35	entropy	5349	0.9574	0.5116	0.8081	0.0010

Results at multiple depths for best criterion via grid search

	Accepted	Denied
Accepted	96	135
Denied	217	1376

Test Data Confusion matrix
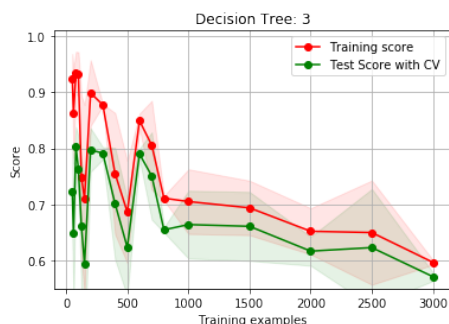
Learning Curve for max_depth = 3
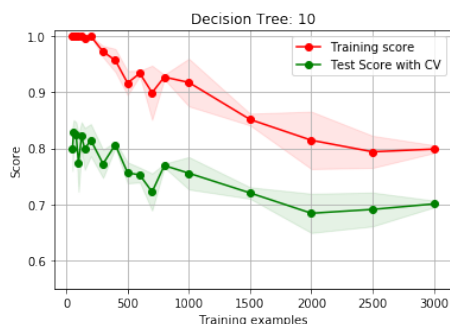
Learning Curve for max_depth = 10

Learning Curve for max_depth = 20

Housing Prices Data

Depth	Criterion	TreeSize	Train %	Train Time	Test %	Test Time
1	gini	3	0.0912	0.0355	0.0890	0.0006
3	entropy	15	0.5816	0.0401	0.5890	0.0003
6	entropy	91	0.6820	0.0442	0.6438	0.0003
10	gini	331	0.8575	0.0566	0.6918	0.0006
15	gini	593	0.9854	0.0695	0.7603	0.0003
20	gini	643	1.0000	0.0559	0.7603	0.0007
25	gini	639	1.0000	0.0550	0.7603	0.0006
35	gini	639	1.0000	0.0570	0.7603	0.0006

Results at multiple depths for best criterion via grid search

	0-1	1-2	2-3	3-4	4-5
0-1	2	6	0	0	0
1-2	4	87	7	0	0
2-3	0	8	15	1	2
3-4	0	0	6	3	0
4-5	0	0	2	0	3

Test Data Confusion matrix (classes in 100ks)
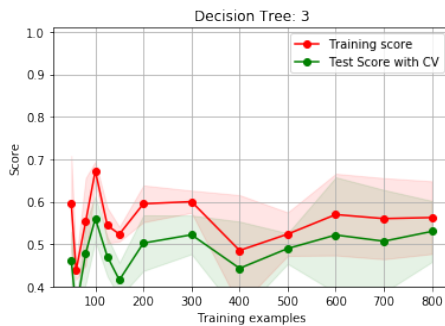
Learning Curve for max_depth = 3   Learning Curve for max_depth = 10   Learning Curve for max_depth = 20

Analysis for Decision Tree

Overall the classifier worked quite well for both datasets, but was extremely prone to overfitting. Examining the results of the decision tree classifier on the two datasets provides numerous observations and basis for analysis, which are provided below.

Effects of dataset size cross validation:

Above, learning curves are provided for both datasets. It is immediately apparent that the size of the dataset greatly affects the performance of the algorithm–though diminishes over time. This makes sense for a few different reasons. The more data we have to train on, the more likely it is that we see the full spectrum of possible variability. Similarly, the broader the set of examples, the less biased our algorithm will be. This is because if we only train on a few data samples, then our algorithm can only make decisions based on the features learned from the small, simple sample size–thus generating a bias (and therefore underfitting).

The learning curves for both datasets clearly level out as dataset size increases, demonstrating less variance with a more informed model. Similarly, cross-validation was used to normalize the training samples and smooth the learnign curve. Without cross-validationg, the model was prone to unrepresentative, biased dataset samples during training and, in-turn, testing.

Tuning Parameters:

The two main paremeters tuned were maximum tree depth and the split criterion. The split criterion measures the quality of a tree split. Both 'gini' and 'entropy' were tested and evaluated using a gridsearch. For the large majority of trials, especially with larger and more complex trees, 'gini' was the more effective splitting criterion–though the results for the two were nearly identical. If anything, Gini was more performant due to it's mathematical simplicity over the entropy formula.

Tree depth provided to be a singificant influencer over the performance of our model–especially for the housing dataset. By allowing for a greater tree depth, we allow for a more complex tree with more possible decisions. A problem emerges when the tree becomes too deep, however. Instead of becoming more robust, the tree begins to overfit with very specific branches for very specific data items. By analyzing the test % vs. depth tradeoff, we can prune unnecessary tree depths for the optimal tree size. This allowed for both a fast, accurate tree with a reasonable footprint.

Performance:

Peformance varied for the two datasets across a few different domains. In terms of runtime, housing prices took longer than the visa data to train (about 10x)–though both were extremely fast on a powerful laptop. This was mainly due to a larger dataset size. As the depth of the tree increased, the training time also took longer for both datasets due to the increased tree size and complexity(as can be seen from the table above).

In terms of accuracy, the visa data started out somewhat high but was able to gain about 4% through optimizing the depth of the tree. This is due to the fact that features like 'salary' and 'job type' proved to heavily influence the approval process. The housing data, however, saw more greater improvements as depth increased. As a more diverse and variable dataset, the model was able to gain a lot of accuracy as the tree grew because it could accommodate more specific feature branches. While the test data started out at <10%, it was able to grow to approximately 76%.

SampleDecision Tree: Permanent Visa with max_depth of 7

2) AdaBoost - Boosted Trees

A boosted decision tree classifer, Adaboost, was the second algorithm applied. It used the same base decision tree as the first model algorithm. Based on the results of the first model, we decided to stick with 'gini' as the criterion for faster training.

US Permanent Visa Data

| Depth	Learning Rate	# Estimators	Train %	Train Time	Test %	Test Time
1	0.1	150	0.8705	24.1484	0.8662	0.0327
3	0.1	100	0.8770	39.7266	0.8679	0.0216
5	0.1	5	0.8761	53.8227	0.8706	0.0021
10	0.1	150	0.9775	82.2192	0.8745	0.0504
15	1	150	0.9775	94.5996	0.8701	0.0420
20	1	150	0.9775	107.0742	0.8618	0.0452

Results at multiple depths for best learning rate/ estimators via grid search

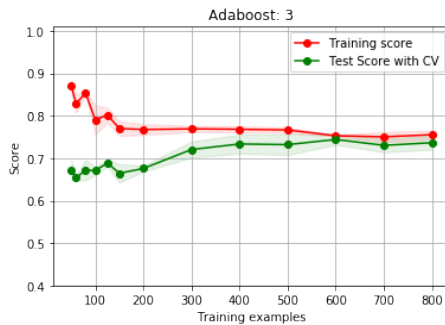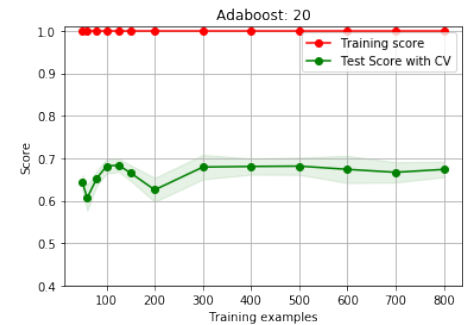Learning Curve for max_depth = 3   Learning Curve for max_depth = 10   Learning Curve for max_depth = 20

Housing Prices Data

Depth	Learning Rate	# Estimators	Train %	Train Time	Test %	Test Time
1	0.1	15	0.6881	6.9413	0.6918	0.0018
3	0.1	5	0.7724	8.2909	0.7945	0.0013
5	0.1	3	0.8291	10.8559	0.7740	0.0011
10	1	50	1.0000	13.6713	0.7945	0.0057
15	1	100	1.0000	8.7526	0.8014	0.0145
20	0.1	15	1.0000	0.6862	0.6781	0.0009

Results at multiple depths for best learning rate/ estimators via grid search

Learning Curve for max_depth = 3   Learning Curve for max_depth = 10   Learning Curve for max_depth = 20

Analysis for AdaBoost Boosted Tree

The AdaBoost boosted tree relates significantly to the original decision tree classifier. It even uses the decision tree as its base classifier. The main difference, however, exists in the boosting. As discussed in class, a boosted tree is essentially an ensemble of weaker trees. This ensemble works to classify with greater accuracy than the original, simplified tree.

Evaluating the performance in the chart and graphs above, it is quickly noticed that the AdaBoost tree works significantly better for both datasets. For visa applications, it is approximately 7% better in the optimal configuration. For housing prices, it i approximately 4% better. Training times take signifcantly longer, which makes sense as the algorithm trains a much larger, ensemble classifier. For example, the larger permanent visa dataset took more than a minute to train for depths greater than 5–compared to less than a second for a normal tree classifier.

In terms of parameters, learning rate, depth, and of estimators were all tuned using a gridsearch to find an optimal model. As of estimators increased, optimal learning rate tended to increase too. For a deeper tree, fewer estimators were needed because the optimal learning rate was achieved sooner. The learning rate, which effects the rate of contribution for each classifier, was optimized between 0.1 and 1, where the learning rate trended higher for deeper trees.

While both types of trees peformed well on the datasets, AdaBoost performed exceptionally well. I believe this is due to the rich set of features and logically classifiable outcomes. The AdaBoosted tree still overclassified at increased depth. By pruning depth to approximately 10 to 15, we were able to achieve an optimal tree.

3) Neural Networks

A neural network (using ScikitLearn's multi-layer perceptron classifier) was the third algorithm applied. A combination of different solvers, learning rates, and scaling was used to observe the functionality of the networks in regards to our dataset.
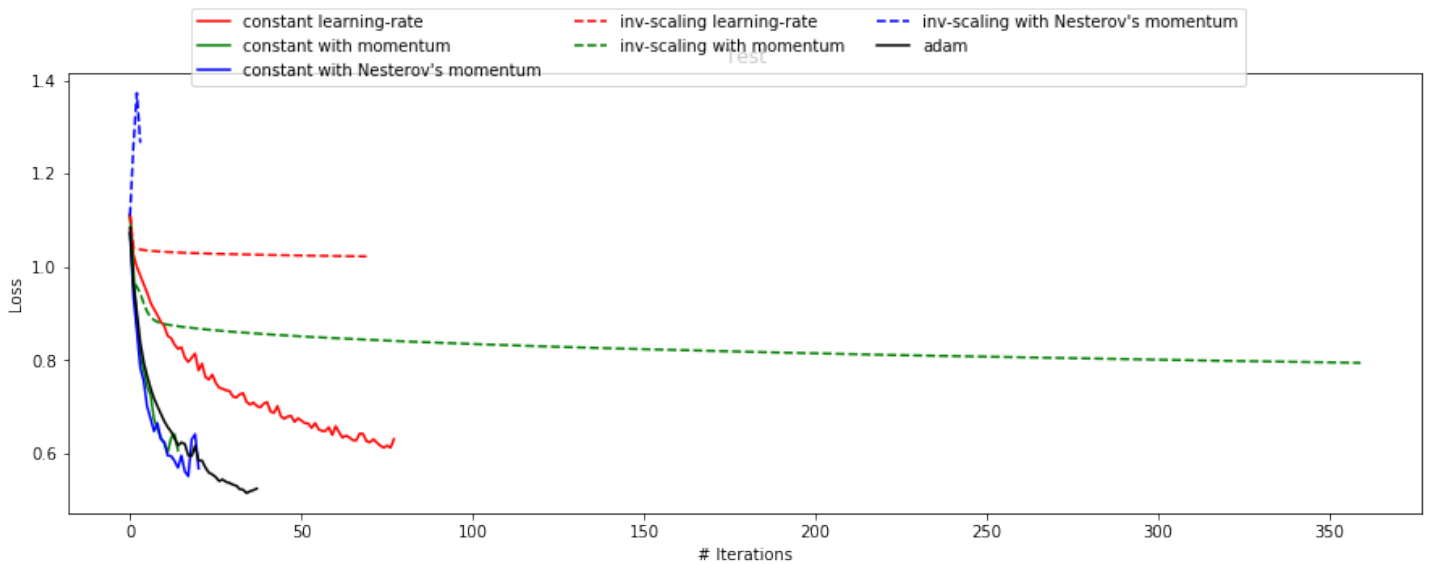
NN: 0.0001/0.001
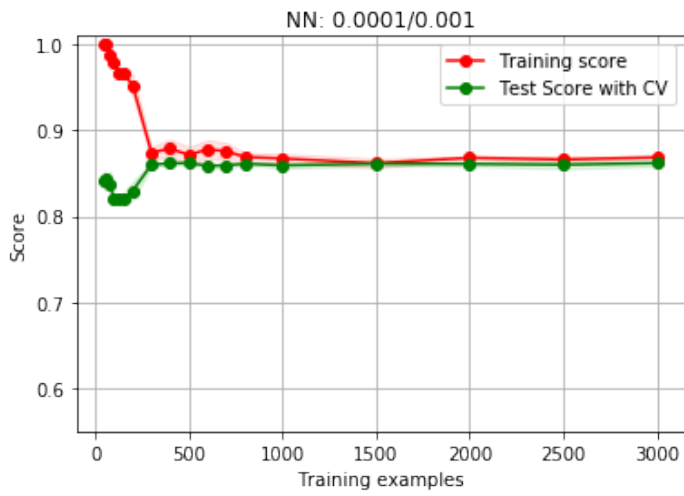
Permanent Visa Applicant NN Learning Curve

NN: 0.0001/0.001
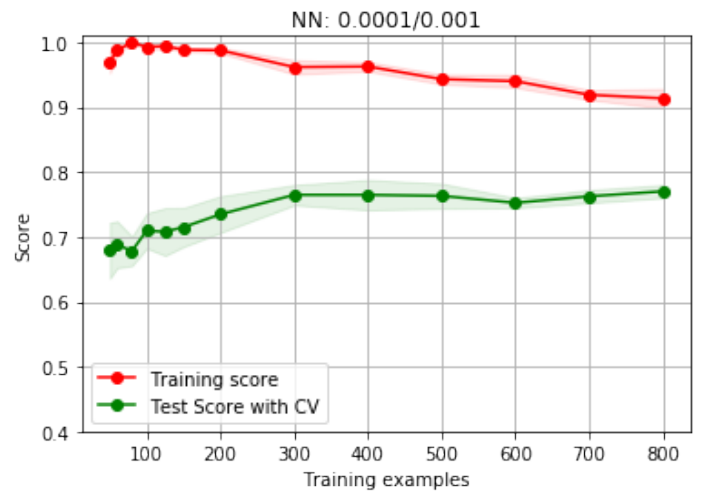
Housing Price NN Learning Curve
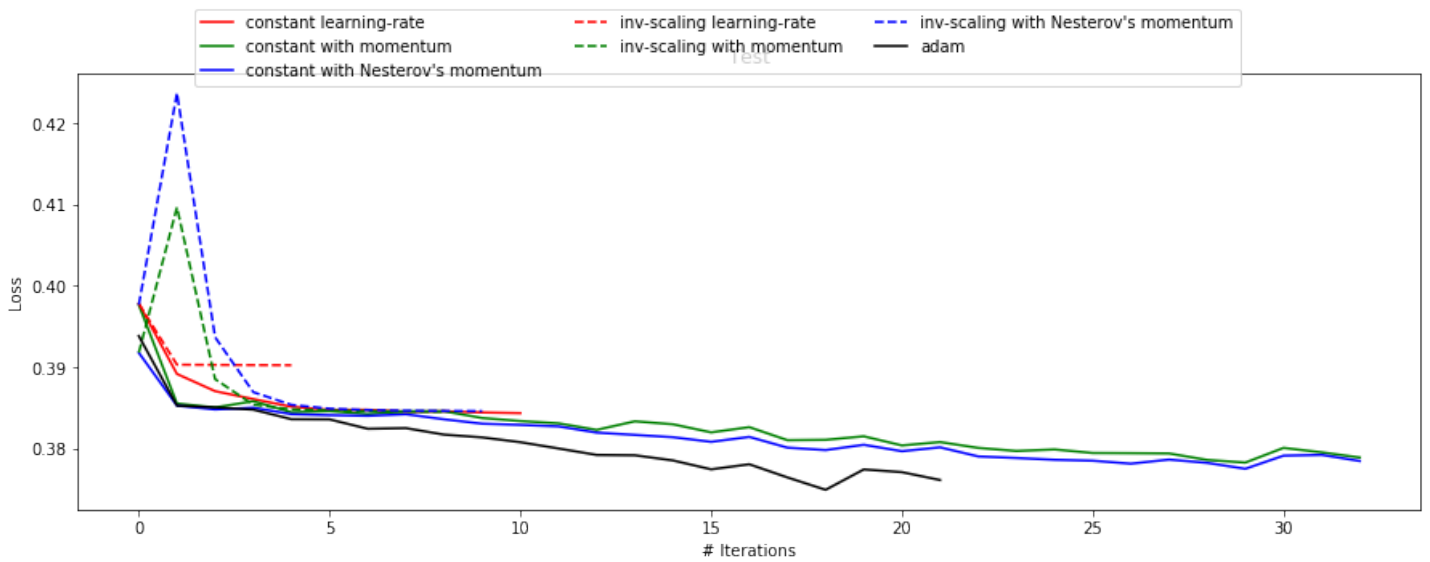
Permanent Visa Applicant NN Configurations Loss Curve

Housing Price NN Configurations Loss Curve

NN Config	Test Score	Loss	Train Time
constant learning-rate	0.8642	0.3842	0.6121
sgd constant learning-rate with momentum	0.8663	0.3823	0.7120
sgd constant learning-rate with Nesterov's momentum	0.8668	0.3779	1.4836
sgd inv-scaling learning-rate	0.8646	0.3905	0.2435
sgd inv-scaling with momentum	0.8646	0.3842	0.3750
sgd inv-scaling with Nesterov's momentum	0.8646	0.3848	0.2740
adam constant learning-rate	0.8677	0.3650	2.1166

Permanent Visa results for various NN configs

NN Config	Test Score	Loss	Train Time
sgd constant learning-rate	0.7664	0.6301	0.4272
sgd constant learning-rate with momentum	0.7498	0.6058	0.0889
sgd constant learning-ratewith Nesterov's momentum	0.7595	0.5669	0.1286
sgd inv-scaling learning-rate	0.6278	1.0219	0.3932
sgd inv-scaling with momentum	0.7009	0.7936	1.9339
sgd inv-scaling with Nesterov's momentum	0.6278	1.2652	0.0205
adam constant learning-rate	0.7795	0.5237	0.2093

Housing Price results for various NN configs

Analysis for Neural Networks

The neural network models, while effective, performed slightly worse than the boosted decision tree. After preliminary exploration, it became apparent that the wide variety of neural network configurations performed differently depending on the given dataset.

In terms of training and testing, cross-validation was used to ensure a robust sampling. The housing price dataset started to converge slowly in regards to the learning curve, whereas the permanent visa dataste convered very rapidly. The permanent visa dataset, by converging rapidly, showed it's low variability. The network model was able to very quickly learn the features–in less than 40–though its feature set is smaller than those of the housing prices. The housing price dataset, conversely, required many more iterations to learn.

Various neural network configurations were used on the test data. Configurations were created by varying the weight optimizer, learning-rate, momentum, and scaling. To aid with speed and peformance, a pre-process Scaler was applied to the datasets. Additionally, a gridsearch approach was applied with various alpha, learning rate, solver, and hidden layer conigurations. While complete results achieved by the various networks are listed above, below we will look at the top performing network for each dataset to try to understand why it was most effective.

For the permanent visa dataset, practically every algorithm performed the same. At first, this perplexed me, as I expected that at least different learning rates and solvers would create different results. It turns out, however, for a dataset that is more easily learnable and converges on fewer iterations, this makes sense. It takes less effort and hyperparemeterization to find functional solutions. To verify that my results were accurate, I also ran a gridsearched neural network with a 2000 iteration limit and various values of solver, learning_rate, and hidden layers.

For the housing price dataset, the 'adam' weight optimizer significantly outperformed the others. This was quite different than the Permanent Visa dataset, which performed rather homogenously. The 'adam' optimizer, or solver, is a 'stochastic gradient-based optimizer.' (From SciKit learn manual) It is noted to work especially well for large datasets. The learning rate was slightly smaller for this configuration, starting with an initial value of 0.01. Since the housing price network took longer to converge and performed quite differently for the various configuration, it is reasonable to believe that with a larger dataset, a more complex and accurate network could be trained. While this would increase training time, the trade off of increased accuracy and lower variance would likely be possible.

4) Support Vector Machines

Support Vector Machines were the fourth algorithm applied.

US Permanent Visa Data

Kernel	Gamma	Train %	Train Time	Test %	Test Time
rbf	0.01	0.8642	20.4459	0.8712	0.1536
rbf	0.05	0.8646	29.0756	0.8712	0.1817
rbf	1.0	0.8752	29.7723	0.8745	0.1805
rbf	2.0	0.8798	28.6184	0.8728	0.1914
poly	0.01	0.8642	9.3467	0.8712	0.0884
poly	0.05	0.8645	20.6110	0.8712	0.0857
poly	1.0	0.8301	24.1879	0.8262	0.1198
poly	2.0	0.8349	19.5531	0.8355	0.0391

Permanent Visa Data

RBF SVM w/ 0.05 gamma

RBF SVM w/ 1.0 gamma

RBF SVM w/ 2.0 gamma

Poly SVM w/ 0.05 gamma

Poly SVM w/ 1.0 gamma

Poly SVM w/ 2.0 gamma

Housing Price Data

| Kernel	Gamma	Train %	Train Time	Test %	Test Time
rbf	0.01	0.7487	0.2301	0.7603	0.0042
rbf	0.05	0.7954	0.2031	0.8219	0.0086
rbf	1.0	0.9226	0.5629	0.7740	0.0042
rbf	2.0	0.9502	0.5996	0.7397	0.0051
poly	0.01	0.6276	0.1581	0.6438	0.0024
poly	0.05	0.7333	0.1733	0.7534	0.0028
poly	1.0	0.9333	0.7254	0.7192	0.0019
poly	2.0	0.9280	1.0244	0.6781	0.0019

Housing Price Data

RBF SVM w/ 0.01 gamma

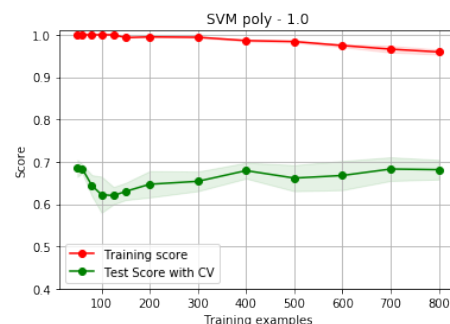
RBF SVM w/ 0.05 gamma

RBF SVM w/ 1.0 gamma

Poly SVM w/ 0.01 gamma

Poly SVM w/ 0.05 gamma

Poly SVM w/ 1.0 gamma

Analysis for Support Vector Machines

The Support Vector Machines peformed well compared to the other algorithms, though, on-average, took significantly longer to train. As there were a lot of parameters to tune, a gridsearch was used yet again along with cross-validation to ensure a robust and effective model. The primary parameters were kernel and gamma.

To begin, two different kernels were used for each dataset: 'rbf' and 'poly'. For both datasets, the rbf kernel proved more effective. The gaussian rbf kernel essentially calculates squared distance between feature vectors and is scaled by gamma. While the poly kernel tends to overfit the training data, rbf does a much better job at securing a robust fit. To tune Gamma, values on a spectrum from 0.01 to 2.0 were used (and results detailed above).

Training time increased signficantly, especially for larger datasets. Since a SVMs require a large amount of iterations and recalculation to converge to an optimal solution, this made sense. It was especially apparent for the larger dataset of permanent visa data. The training time took approximately 30 seconds for the various configurations where as the smaller housing price dataset took less than a second on average.

From the learning curves, we see that a highly succesful training rate leads to greater overfitting. This makes sense, as the SVM becomes too rigid to the input data, and is not robust. Similarly, when the kernel has too low of a gamma and does not accomodate feature relations properly, the model underfits and performs poorly on both training and test data.

5) K-nearest Neighbors

K-nearest Neighbors was the fifth algorithm applied.

Permanent Visa Data

Kernel	K	Train %	Train Time	Test %	Test Time
uniform	1	0.9708	0.3422	0.8158	0.0080
uniform	2	0.8978	0.3619	0.7708	0.0094
uniform	3	0.8989	0.3603	0.8394	0.0106
uniform	4	0.8854	0.3836	0.8202	0.0103
uniform	5	0.8892	0.3996	0.8503	0.0107
uniform	10	0.8742	0.4736	0.8525	0.0126
uniform	15	0.8739	0.5735	0.8640	0.0134
uniform	20	0.8738	0.6083	0.8673	0.0168
uniform	30	0.8696	0.8306	0.8662	0.0219
uniform	50	0.8686	1.1392	0.8618	0.0315
distance	1	0.9708	0.2706	0.8158	0.0071
distance	2	0.9675	0.4111	0.8103	0.0113
distance	3	0.9743	0.4050	0.8279	0.0096
distance	4	0.9741	0.3766	0.8311	0.0092
distance	5	0.9756	0.4084	0.8377	0.0157
distance	10	0.9762	0.4899	0.8470	0.0142
distance	15	0.9767	0.6011	0.8509	0.0169
distance	20	0.9782	0.6799	0.8591	0.0290
distance	30	0.9782	0.9327	0.8624	0.0255
distance	50	0.9782	1.2837	0.8635	0.0397

Permanent Visa Data

KNN Comparision
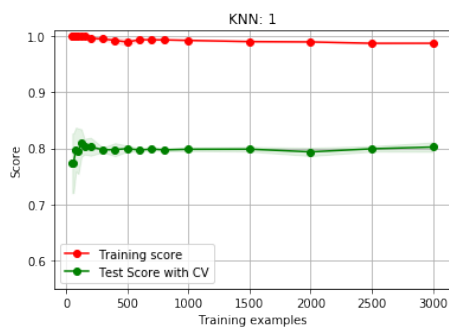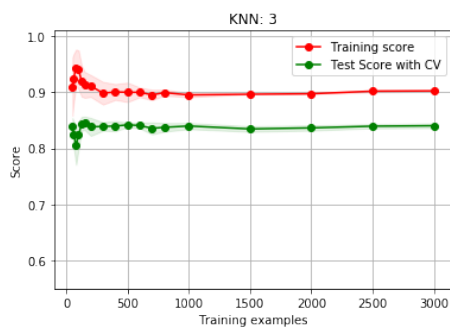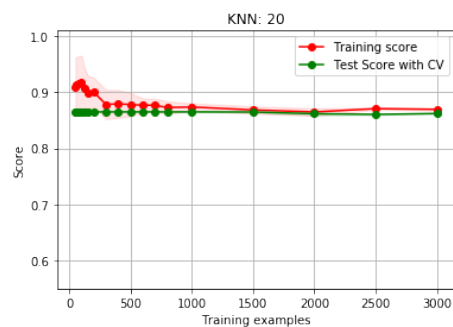

Uniform KNN - 1

Uniform KNN - 3

Uniform KNN - 20



Distance KNN - 1

Distance KNN - 3

Distance KNN - 20

Housing Price Data

| Kernel	K	Train %	Train Time	Test %	Test Time
uniform	1	1.0000	0.0343	0.6507	0.0015
uniform	2	0.8222	0.0304	0.6644	0.0011
uniform	3	0.8130	0.0283	0.7055	0.0011
uniform	4	0.7824	0.0379	0.6986	0.0012
uniform	5	0.7602	0.0351	0.7260	0.0012
uniform	10	0.7341	0.0324	0.7055	0.0019
uniform	15	0.7218	0.0369	0.7192	0.0016
uniform	20	0.7134	0.0466	0.7329	0.0013
uniform	30	0.7004	0.0488	0.7055	0.0022
uniform	50	0.6812	0.0635	0.6986	0.0020
distance	1	1.0000	0.0244	0.6507	0.0009
distance	2	1.0000	0.0261	0.6575	0.0010
distance	3	1.0000	0.0293	0.7055	0.0010
distance	4	1.0000	0.0282	0.6986	0.0010
distance	5	1.0000	0.0346	0.7397	0.0011
distance	10	1.0000	0.0356	0.7123	0.0012
distance	15	1.0000	0.0385	0.7192	0.0017
distance	20	1.0000	0.0468	0.7534	0.0021
distance	30	1.0000	0.0501	0.7260	0.0016
distance	50	1.0000	0.0742	0.7055	0.0021

KNN Comparision

Housing Price Data

Uniform KNN - 1


Uniform KNN - 3

Uniform KNN - 20

Distance KNN - 1

Distance KNN - 3

Distance KNN - 20

Analysis for k-Nearest Neighbors