

CS 7641 Machine Learning

Assignment 3

Philip Bale
pbale3

Due Sunday April 1st, 2018 11:59pm

Introduction

This assignment explores unsupervised learning and dimensionality reduction. It begins by examining clustering algorithms, specifically k-means and expectation maximization. It then proceeds to cover four dimensionality reduction algorithms: principal components analysis, individual components analysis, randomized projections, and random forests. After running these six algorithms on the original datasets and observing the results, the results are then piped into a neural network learner for further examination.

Datasets

The datasets chosen were the same datasets chosen for assignment 1. The first dataset is the US permanent visa dataset. This dataset is interesting due to its potential to aid in the visa application process from a cost and time savings potential. It could also enable confidence in those interested in applying for a US permanent visa but doubting their chances of acceptance. At the end of the day, the goal is it to try to determine the application result before time, money, and other resources are spent.

The second dataset is a home sale price prediction dataset taken from an ongoing Kaggle competition. This dataset is interesting for two primary reasons: real-world applicability and participating in a Kaggle challenge. First, modeling home prices is both a difficult and lucrative task. If one can successfully model home sale prices on large sets of data, he/she can make large amounts of money investing in real estate when he/she detects outliers in listed price vs. what it is expected to sell for. This applies to flipping, investing, and remodeling. Second, the dataset is part of an ongoing Kaggle competition that does not have a winning solution yet. By taking part of the competition, the dataset presents the opportunity to work towards a winning solution and advance one's algorithms over time.

Part 1: Clustering Algorithms




Introduction

K-means clustering is the first algorithm applied to the datasets and expectation maximization is the second. Both algorithms work by clustering: gathering groups of instances together based upon their features. The rationale is that similar instances will likely be labeled the same way—such as identical visa applications obtaining the same outcome.



1) k-means clustering

Overview

K-means works by clustering n instances into k -clusters of similarity using least-squares Euclidean distance between the instances. In practice, the algorithm converges on 'mean' for each cluster that is representative of the members of that cluster. A variety of cluster sizes were tested to find the best parameters possible.



Perm Visa Sum of Square Errors for Clus- Perm Visa Log Likelihood vs. # Compo- Perm Visa Scoring for k-means and ex-
ters vs. # Clusters nents pectation maximization



Housing Sum of Square Errors for Clus- Housing Log Likelihood vs. # Compo- Housing Scoring for k-means and expec-
ters vs. # Clusters nents tation maximization

k-Means Analysis

Observing the graphs above, it is clear to see that varying the number of clusters used has a clear impact on the performance of k-Means clustering. For both datasets, as the number of clusters increases, the clusters are more able to represent the data. Logically, there are instances where increasing number of clusters will decrease the accuracy instead of increasing it, such as when just starting out and before convergence (using Euclidean distance).

The first measurement used to determine the effectiveness is the sum of square errors (SSE). The SSE measures how far away an instance data point is from the mean of its cluster. As the number of clusters increases, the SSE noticeably drops and then converges. This makes sense because at a certain point, adding more clusters is overfitting and not necessary to get the all training data into its best possible fit.

From the scoring data, it is shown that the permanent visa data, performs remarkably well with a small number of clusters and does not show any noticeable improvement by increasing clusters. This is due to the fact that the permanent visa data is extremely homogenous and does not contain many outliers at all. On the other hand, the housing price data is much more susceptible to changes in number of clusters. As the number of clusters increases, the testing data gradually increases in accuracy before leveling off. Since the housing data is much more varied and complex, there are intricacies of the data that require more clusters to capture well.

| Clusters | 2 | 5 | 10 | 15 | 20 | 25 | 30 | 35 | 40 | 50 | 60 | 70 |
|-----------------------|--------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| PERM VISA | | | | | | | | | | | | |
| SSE | 108717 | 71834 | 47453 | 37701 | 31090 | 27611 | 25517 | 23874 | 22410 | 20267 | 18532 | 17331 |
| Log Likelihood | -9.44 | 2.67 | 9.57 | 13.25 | 13.90 | 16.01 | 15.78 | 15.29 | 17.56 | 19.12 | 20.04 | 20.47 |
| k-Means AMI | 0.022 | 0.008 | 0.005 | 0.005 | 0.004 | 0.005 | 0.004 | 0.004 | 0.004 | 0.004 | 0.004 | 0.004 |
| k-Means ACC | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 |
| EM AMI | 0.022 | 0.007 | 0.005 | 0.005 | 0.006 | 0.005 | 0.006 | 0.007 | 0.006 | 0.006 | 0.007 | 0.007 |
| EM ACC | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 | 0.865 |
| HOUSING | | | | | | | | | | | | |
| SSE | 14840 | 10217 | 8265 | 7256 | 6589 | 6106 | 5687 | 5339 | 5063 | 4589 | 4276 | 4024 |
| Log Likelihood | | | | | | | | | | | | |
| k-Means AMI | 0.105 | 0.134 | 0.120 | 0.103 | 0.091 | 0.080 | 0.086 | 0.090 | 0.086 | 0.081 | 0.083 | 0.081 |
| k-Means ACC | 0.628 | 0.634 | 0.695 | 0.702 | 0.694 | 0.688 | 0.695 | 0.704 | 0.705 | 0.715 | 0.719 | 0.726 |
| EM AMI | 0.010 | 0.065 | 0.073 | 0.073 | 0.078 | 0.073 | 0.081 | 0.085 | 0.077 | 0.078 | 0.065 | 0.064 |
| EM ACC | 0.628 | 0.631 | 0.631 | 0.644 | 0.657 | 0.666 | 0.682 | 0.688 | 0.686 | 0.677 | 0.673 | 0.679 |

Table of Housing Data Results for Cluster

2) Expectation Maximization

Overview

Expectation Maximization is the second algorithm applied to the datasets and, similar to k-means, is a clustering algorithm. Expectation Maximization works by iteratively finding the maximum likelihood of parameters leading to a labeling of an instance despite possibly not having all data or parameters. For our examples, we used Scikit-learn's Gaussian mixture models to implement the Expectation Maximization algorithm. A varying number of mixture components (or number of distributions) were used to determine the best possible parameters for the clustering.

Expectation Maximization Analysis

Expectation maximization performed only slightly worse than k-means on the datasets. Instead of using a sum of square errors calculation, a log likelihood is calculated to effectively determine the probability of successful labeling. Interestingly, the housing dataset converges quite quickly to a near-peak log likelihood where as the permanent visa dataset takes a bit longer. This makes sense, as the permanent visa dataset is much larger and while an indicator of classification performance and determining factor for component count, it does not guarantee how well the algorithm will perform using such settings.

In terms of scoring, while k-means performed slightly better, it isn't by much for the housing dataset—and it was insignificantly better for the permanent visa dataset. The adjusted mutual info score, which helps to determine the differences between clusters while accounting for chance, also performs similarly for expectation maximization compared to k-means. Overall, while k-means performed better in our trials, it is reasonable to believe datasets exist that would fare better using expectation maximization.

Part 2: Dimensionality Reduction Algorithms

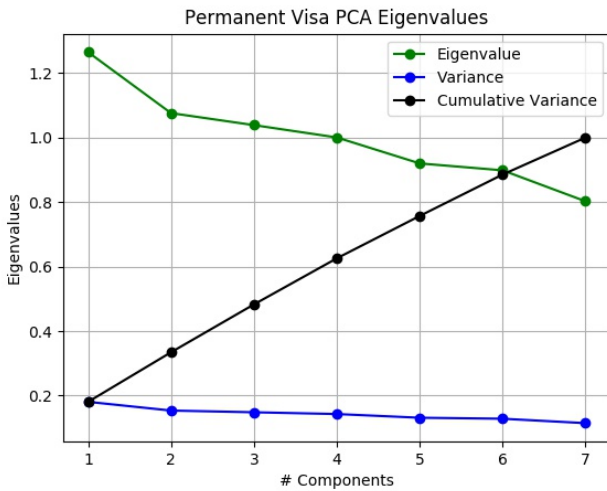
Introduction

Part 2 deals with dimensionality reduction algorithms. The four algorithms used are principal components analysis, individual components analysis, randomized projections, and random forests. After running the algorithms on both datasets, an analysis is provided on the results.

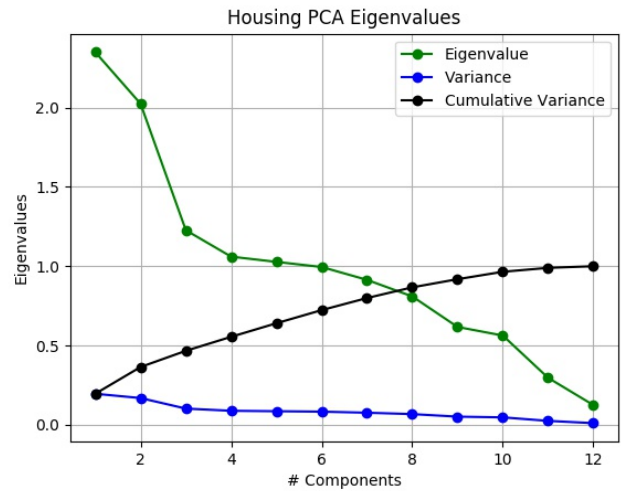
1) Principal Components Analysis (PCA)

Overview

The first dimensionality reduction algorithm, Principal component analysis is a statistics approach to finding vectors that maximize variance and thus help to determine components that are correlated. Each subsequent component is found with the intent to be orthogonal to the preceding component. The resulting eigenvalue matrix from PCA is therefore maximized for covariance.



Permanent Visa Principal Components Analysis



Housing Principal Components Analysis

Analysis


Principal component analysis seeks to reduce the number of dimensions in the data without sacrificing data quality. The principal component analysis results show that both the eigenvalues and the variance decrease for both datasets as the number of components is increased. For the permanent visa dataset, the size of the eigenvalues starts to level off, indicating that there are features that are potentially unnecessary and removable. The housing dataset, however, has more of a continuous downwards trend for its eigenvalues, indicating that removing features may not be a wise thing to do.

The variance graphs also provide interesting views into the ability of PCA to reduce dimensionality without sacrificing data quality. While the housing dataset indicates a higher variance between different features, the permanent visa dataset proves to be more evenly distributed. Since we are trying to maximize variance between the different components (so that we most accurately represent the higher dimension data), we want to choose a number of components that demonstrates such. For the permanent visa dataset, that number appears to be around 5 components and for the housing dataset it appears to be around 9 components.


2) Independent Components Analysis (ICA)

Overview

The second dimensionality reduction algorithm, independent components analysis, is an approach to separating a mixture of a data into appropriate subcomponents. As discussed in lecture, a good example of what ICA is used for is the cocktail problem; where one needs to separate various sounds into their sources: a tv show, humans, car noises, etc. Kurtosis is used as a measurement of how gaussian the derived components are.



Permanent Visa Independent Components Analysis



Housing Independent Components Analysis


Analysis

While PCA sought to maximize variance, ICA seeks to separate mixed data into subcomponents. As a dimensionality aglorithm, independent components analysis also wants to minimize dimensionality while preserving data quality. Using the kurtosis measurement, we are able to measure the spikiness of the data distribution. It's important to note that kurtosis is sensitive to outliers and therefore not always robust to measuring gaussianity. Similar to PCA, it is observed that the permanent visa dataset is significantly more homogenous then the housing dataset.


3) Randomized Projections

Overview

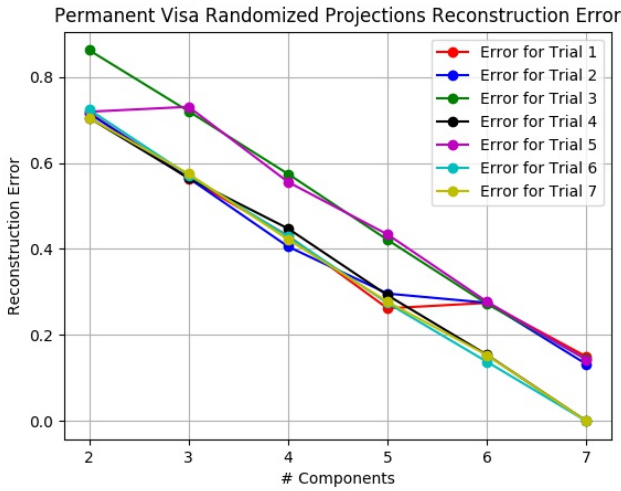
The third dimensionality reduction algorithm, randomized projections, is an approach that randomly generates a projection matrix that attempts to create a lower dimension representation of the data that is approximately accurate to its original state. By varying the number of components to project, we can run varoius tests on how well the lower dimension data captures the original.




Permanent Visa Randomized Projections Pairwise Correlation



Housing Randomized Projections Pairwise Correlation



Permanent Visa Randomized Projections Reconstruction Error



Housing Randomized Projections Reconstruction Error

Analysis


Randomized projections, of all the dimensionality reduction algorithms, was the most susceptible to variation in performance due to its random nature. In such, various trials were run, each varying the random state and maintaining that same random state for a variety of number of components. The first measure used to determine the appropriate number of dimensions to reduce to was pairwise distance between the original and reduced data. From the graphs, it is easy to see that the distance (akin to difference in the instances) converges around 5-6 components for the permaentn visa dataset and 9-10 for the housing dataset.

The second, measurement used was reconstruction error. For both datasets, the reconstruction error decreases significantly as the number of dimensions is increased. This makes sense as the more dimensions available, the more easily the initial data can be reconstructed.

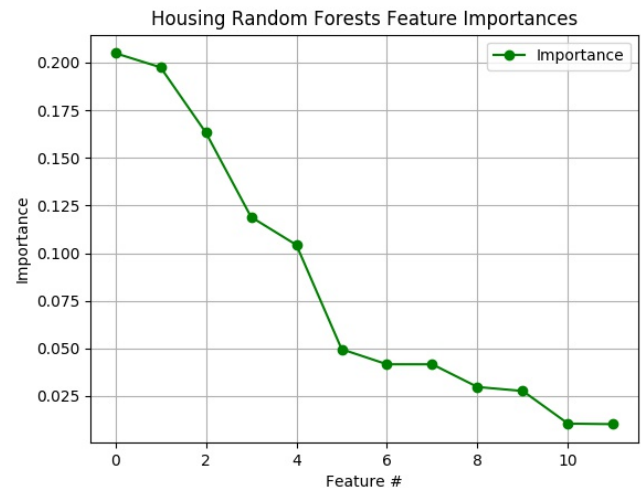
4) Random Forest Feature Selection

Overview

The fourth, and last, dimensionality reduction algorithm, random forest feature selection, is an approach that uses an ensemble of decision trees conditioned on different features. By training the decision tree and observing the impact of each feature by its ability to classify data correctly, we can select the most important features and disregard unimportant features.



Permanent Visa Random Forest Feature Importances (Descending Order)



Housing Random Forest Feature Importances (Descending Order)

Analysis


IText

Part 3: Dimensionality Reduction and Clustering


Overview

In this section, clustering algorithms are run on the results of the dimensionality reduction algorithms and then compared. All dimensionality reduction and all clustering algorithms from above are used.


k-Means after Dimensionality Reduction




Perm Visa SSE



Perm Visa Scoring



Housing SSE




Housing Scoring


Analysis

Text


Expectation Maximization after Dimensionality Reduction




Perm Visa Log Likelihood



Perm Visa Scoring



Housing Log Likelihood



Housing Scoring


Analysis

Text

Part 4/5: Dimensionality Reduction, Clustering, and Neural Networks

Overview

In this section, similar to part 3, neural networks are run on the results of the dimensionality reduction algorithms and the clustering algorithms, and then compared.



Housing NN after dimensionality reduction



Housing NN after clustering

Dimensionality Reduction + NN Analysis

Clustering + NN Analysis

Conclusion

Todo conclusion