CS 7641 Machine Learning
Assignment 2

Philip Bale
pbale3

Due Sunday March 11th, 2018 11:59pm

Part 1: Neural Network Optimization

Introduction

Part 1 of the assignment surrounds using randomized optimization to find the best possible weights for a specific neural network. In assignment 1, backpropogation was used to find optimal parameters for a neural network. This neural network took in various input features for US permanent visa applicants and then attempted to predict the outcome of an application. After various tests, I found the optimal parameters of: 6-node input layer, one hidden layer with 100 nodes, one output node, and about 500 iterations.

For assignment 2, three other optimization strategies are employed to find optimal weights: randomized hill climbing, simulated annealing, and genetic algorithms. The details of each are discussed in-depth further in this section and an analysis is performed for each strategy to evaluate its results. It is important to note that this network does tend to converge quite quickly due to its size and consistant distribution.

I primarily chose this problem domain because, as someone who has worked with a large number of first-generation visa holders and immigrants, I am extremely interested in building tools to help others to achieve the same. At the end of the day, the goal is it to try to determine the application result before time, money, and other resources are spent.

1) Backpropogation (Assignment 1)

Overview

The first weight-finding algorithm used was backpropogation. Backpropogation works by essentially calculating the error at the end of a network, and then working backwards to minimize that error over various iterations. An error (or loss) function is effectively minimized over time using this backpropogation technique. As discussed in assignment 1, the permanent visa is rather large and robust. It is quickly learnable by various different learners and in such backpropogation found significant success.
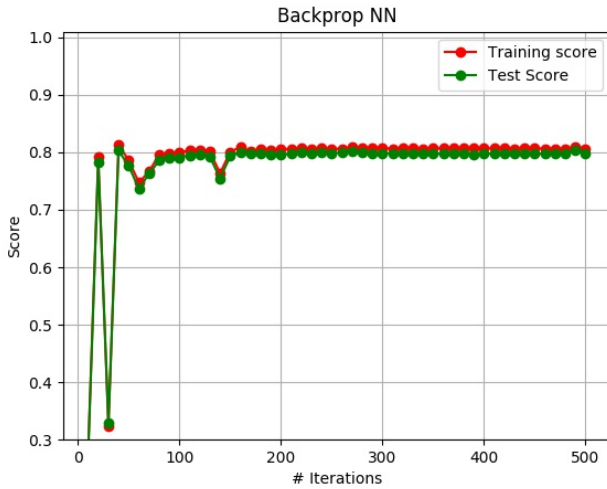
Backprop NN Success Rate vs. Iterations
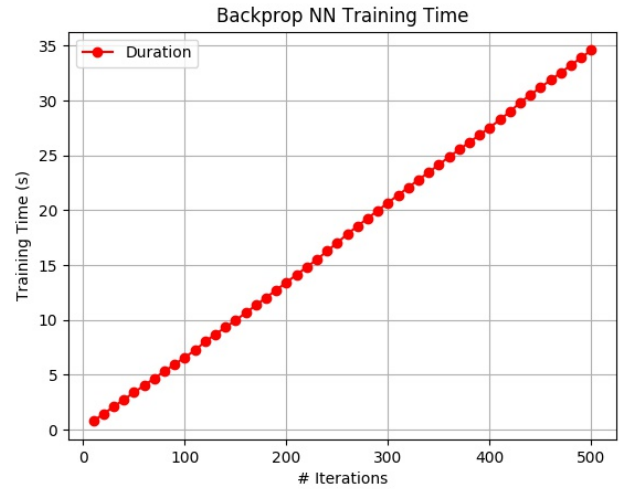

Backprop NN Training Time

Right around 50 iterations, the network begins to converge at around an 80% successs rate. Seeing as the training and test score track either rather closely, it is apparent that the dataset is rather robust and consistent. One thing to note is that the training time scales linearly with the number of iterations–which makes sense since the same amount of calculations with similar complexity are performed on each iteration of backpropogation.

2) Randomized Hill Climbiing

Overview

The second weight-finding algorithm used was randomized hill climbing. Randomized hill climbing works by taking a random starting point and then incrementally attempting to improve on that point. In the context of a neural network trying to find weights, randomized hill climbing selects random weights and then moves in a direction so as to try to find a better result for that weight–akin to trying to move up an optimization 'success hill'.

One thing to note is that we are using randomized hill climbing, not random restart hill climbing. In such, the algorithm is prone to getting caught in local optimizations, or local maximums.

Randomized Hill Climbing NN Success Rate vs. Iterations

Randomized Hill Climging NN Training Time

High accuracy results are achieved right around 125 iterations+. While this is subject to randomness, by looking at the results it is shown to become rather consistant. By looking at the score results, one can see various instances of the randomized hill climbing getting caught in local optimas and being unable to escape. Such is the case around 220, 320, and 350 iterations.

Similar to backpropogation, training time scales linearly with the number of iterations run. Training time tends to be a bit faster using randomized hill climbing because of a reducation in calculations necessary. Whereas backpropogation needed to do calcluations to minimize error moving backwards through the network, randomized hill climbing simply needs to move in one direction and determine if the new weights are better.
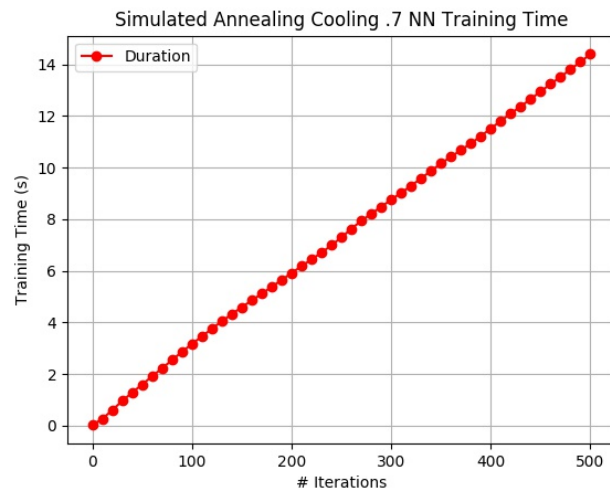
3) Simulated Annealing

Overview

The third weight-finding algorithm used was simulated annealing. Simulated annealing works by taking a random solution and then samples nearby alternatives. By comparing the alternatives to the original solution, the optimizer decides to either stick with the original solution or move to the new one. Using a temperature and cooling parameter, the algorithm is more open to worse solutions at first but gradually moves towards only accepting better solutions.

Various coolng parameters were used to try to determine the best simulate annealing approach. Below, success rates by number of iterations are showed for simulated annealing approaches with cooling parameters of .15, .35, and .7. The parameter of .15 and .35 take significantly longer to converge to the optimal solution than did a higher parameter. While part of this can be attribute to luck (the algorithm could have randomly found itself in an optimal solution earlier on)–part of this too is the way the algorithm operates by its parameters. The .7 cooling temperature here proved most effective.

Simulated Annealing NN Success Rate vs. Iterations w/ .15 Cooling

Simulated Annealing NN Success Rate vs. Iterations w/ .35 Cooling

Simulated Annealing NN Success Rate vs. Iterations w/ .7 Cooling

Simulated Annealing NN Training Time w/ .7 Cooling

4) Genetic Algorithms

Overview

The fourth weight-finding algorithm used was a genetic algorithm. Genetic algorithms works by starting with an initial solution and then making modifications in an attempt to improve the solution. The modifications generally allowed are mutation (changing random parts of the solution), crossover (taking specific sections from various solutions and combining them), and selection (selecting certain sections from a solution to use again). In the context of a neural network, we can use genetic algorithms to make modifications to our network's weights.

In our testing, a population size of 50 was used in order to get a diverse initial sampling of possible solutions. The number of instances to mate (aka crossover) and to mutate was varied throughout the trials.

Genetic Algorithm Success Rate v. Iterations w/ 10 Mate, 10 Mutate, 50 population

Genetic Algorithm Success Rate v. Iterations w/ 10 Mate, 20 Mutate, 50 population

Genetic Algorithm Success Rate v. Iterations w/ 20 Mate, 10 Mutate, 50 population

Genetic Algorithm Success Rate v. Iterations w/ 20 Mate, 20 Mutate, 50 population

Trials were run with the following mate/mutate combinations: 10/10, 10/20, /20/10, and 20/20. All trials were run with a population size of 50. As can be observed from the graphs, both mating and mutating appeared to be highly effective in finding the optimal solution. In fact, the optimal solution converged quite quickly–in under 30 iterations for each trial. This is partly due to the fact that the data is easily learnable and does not appear to maintain very many trapping local maximums.

Training times scaled roughly linearly with the number of iterations ran. This makese sense because, from a performance perspective, the number of mutation calculations from iteration to iteration is more or less the same.

Genetic Algorithm Training Time w/ 10 Mate, 10 Mutate, 50 population

Genetic Algorithm Training Time w/ 20 Mate, 20 Mutate, 50 population

Conclusion

The optimization algorithms used all inevitably produced similar results. What differed, however, was the amount of training time required, the tuning of parameters, and how many iterations were required to achieve a consisent, optimal solution. All-in-all, the genetic algorithm approach consistantly produced the best results for our network. This makes sense, as the data was rather homogenous and there appeared to be very few outliers. By learning the training data well, the learner was able to perform similarly well on the (nearly identical) test data.

Part 2: Optimization Problem Domains

Introduction

For this part of the assignment, three different optimization problems are examined: continous peaks, the traveling salesman, and flip flop. The three optimization techniques from above are used (randomized hill climbing, simulated annealing, and genetic algorithms) as well as another algorithm, MIMIC.

MIMIC, similar to the other algorithms, works to find the globally optimal solution. Unlike the other algorithms, however, it retains knowledge of previous iterations and uses this information to more efficiently find better solutions. MIMIC is particularly strong in regards to problems that maintain patterns between subsets of thier parameters.

Below, we apply each of the optimization algorithms to each of the three optimization problems. We then look at how each optimization problem's characteristics may favor one algorithm to another, as well as evaluate the efficiency of each.

1) Continous Peaks

Overview

The continous peaks problem is an extension of the four peaks problem that allows for a wide variety of local maximums. This algorithm is particularly interesting due to the potentially large number of local maximums. It is especially difficult for an algorithm that cannot escape local peaks to perform well.

Fitness vs Iterations of 4 Optimization Strategies
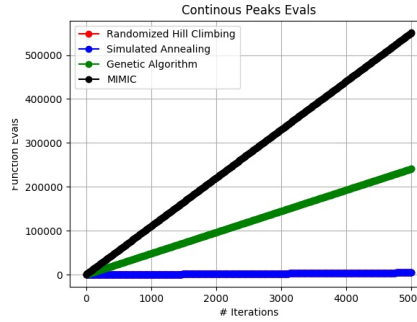
Function Evals vs Iterations of 4 Optimization Strategies

Training Time vs Iterations of 4 Optimization Strategies

Analysis

In order to ensure accuracy and a fair distribution, 5 different trials were used of 5000 iterations each. Each optimization strategy was also allowed to evaluate over a parameter range in an attempt to find the optimal parameters for that strategy.

For randomized hill climbing, there were no other parameters to tune. The algoritm was simply run for various iterations (multiples of 10) in a set of 5 trials. In terms of fitness function performance, it actually performed quite well and required no significant function evaluations. It was similarly fast to train due to the extreme simplicity of the algorithm itself.

Iterations	Fitness	Training Time	Function Evals
0	16.0	5.5038e-05	11
10	16.0	7.64649999999e-05	21
50	22.0	0.000212433	61
100	22.0	0.000370814	111
200	23.0	0.000717487	211
300	23.0	0.001094822	311
500	26.0	0.001666152	511
750	31.0	0.002682947	761
1000	51.0	0.003609291	1011
1500	62.0	0.005517988	1511
2000	74.0	0.007139624	2011
2500	84.0	0.008715217	2511
3000	90.0	0.010594833	3011
3500	95.0	0.012718942	3511
4000	99.0	0.014618306	4011
5000	99.0	0.018770859	5011

Simulated Annealing Results w/ .15 Cooling

Simulated annealing was slightly more complex, where cooling efficients of 0.15, 0.35, 0.55, 0.75, and 0.95 were used. It was found that the efficient 0.15 was the most effective here. This low cooling rate was effective because it helped to slowly overcome any local maximums but not work too aggressively to skip any optimal solution. By iteration 3k, Simulated Annealing took a clear lead over the other solutions in terms of fitness performance. Similar to randomized hill climbing, its number of function evals and training time was extremely low.

The genetic algorithm employed used a population size of 100 and various configurations of 50, 30, and 10 mutations/mate combinations. The graph above depicts the succes rates when using 30 mutation and 30 matings. As can be seen, it's performance was rather simular to randomized hill climbing–which roughly makes sense due to the random nature of both algorithms. The number of function evaluations was higher for the genetic algorithm due to the amount of attempts at mutating and mating–whereas the training time was practically irrelevant.

Mimic's parameters were set to 100 samples and a 'to-keep' number of 50. The threshold for the discrete dependency tree mimic used was sampled from the range 0.1, 0.3, 0.5, 0.7, and 0.9. While MIMIC started off relatively string in terms of

fitness performance, it thrashed around the 60-80 range for the most part. I had suspected that MIMIC would have difficulty on the continous peaks due to the large amount of randomness in the data. The number of functional evaluations required for MIMIC was significnatly higher than the other algorithms, as was its training time. This was primarily due to the fact that MIMIC is signifciantly more complex to operate than the other algorithms.

Overall, Simulated Annealing proved to be the best strategy for the continous peaks problem, largely due to it's ability to handle large amounts of randomness well and to avoid getting caught in such local maximums. It's low training time and number of function evaluations further aided its case.

2) Traveling Salesman

Overview

The traveling salesman problem is a frequently used, classic problem that focuses on a salesperson trying to minimize their round-trip distance between any number of cities. This problem is particularly interesting due to it's real-world implications, such as route planing for delivery trucks and everyday errands–it also has no known polynomial time solution!
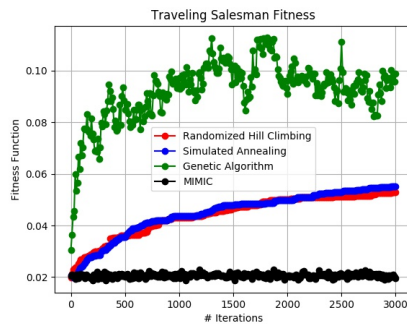
Fitness vs Iterations of 4 Optimization Strategies
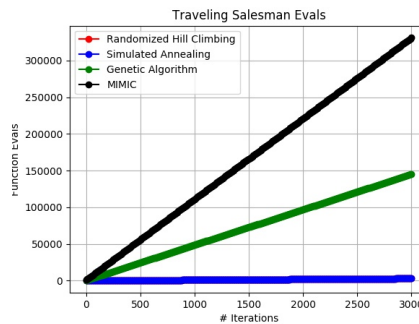
Function Evals vs Iterations of 4 Optimization Strategies
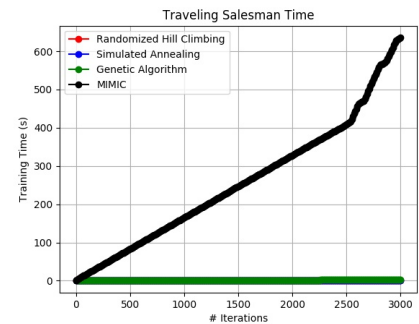
Training Time vs Iterations of 4 Optimization Strategies

Analysis

For this optimization problem, in order to ensure accuracy and a fair distribution, 5 different trials were used of 3000 iterations each. 50 Random points were generated and used as the destinations for our faux salesman. Each optimization strategy was also allowed to evaluate over a parameter range in an attempt to find the optimal parameters for that strategy.

For randomized hill climbing, there were no other parameters to tune. The algoritm was simply run for various iterations in a set of 5 trials. In terms of fitness function performance, it performed mediocrely and required no significant function evaluations. It was similarly fast to train due to the extreme simplicity of the algorithm itself. The reason for the stragegy's mediocrity can easily be attributed to the fact that random guessing does not perform well when trying to develop the most optimal path due to the large number of possibilities and low number of optimums.

Similar to the last problem, simulated annealing was slightly more complex, where cooling efficients of 0.15, 0.35, 0.55, 0.75, and 0.95 were used. It was found that the efficient 0.55 was the most effective here. Interestingly enough, the simulated annealing approach very closely tracked the randomized hill climbing results. This was pretty cool, and clearly demonstrated the effect of the problem on the ability of a solution to perform well. As both were prone to picking random solutions and making marginal improvements on those solutions, they shared the same strenghts and weaknesses for trying to plan an optimal path. Similar to randomized hill climbing, simulated annealings' number of function evals and training time was extremely low.

Iterations	Fitness	Training Time	Function Evals
0	0.0305500487733	0.011139866001	573
50	0.0535358597005	0.0619086620027	2984
100	0.0670653896105	0.106549986003	5404
200	0.0754973396223	0.190760427002	10194
300	0.0843530928911	0.276369227002	15021
500	0.0847430696768	0.430181398007	24690
750	0.0824751535861	0.62021527501	36716
1000	0.097111513149	0.799495104007	48778
1500	0.100315033444	1.153653522	72921
2000	0.0908447200981	1.51476069201	97068
2500	0.111066212921	1.875438519	121202
3000	0.0986916191186	2.27946003001	145282

Genetic Algorithm Results w/ population 100, 30 mutations, 30 matings

The genetic algorithm employed used a population size of 100 and various configurations of 50, 30, and 10 mutations/mate combinations. The graph above depicts the succes rates when using 30 mutation and 30 matings. As can be seen, it's performance was exceeding better than the other algorithms. The mutations and matings were able to successfully identify the most optimal paths and plan extremely well for the hypothetical salesman. This can be attributed to the fact that traveling salesman can be broken into various subjourneys that the genetic algorithm can learn and then attempt to combine for an optimal journey. The number of function evaluations was higher for the genetic algorithm due to the amount of attempts at mutating and mating–whereas the training time was practically irrelevant. The genetic algorithm was the clear winner in terms of performance for the TSP.

Mimic's parameters were set to 100 samples and a 'to-keep' number of 50. The threshold for the discrete dependency tree mimic used was sampled from the range 0.1, 0.3, 0.5, 0.7, and 0.9. MIMIC performed rather poorly on this problem. This was very likely due to the fact that pattern finding is of almost no use in this problem domain and that previous knowledge also is mostly unnecessary when trying to find optimal paths. The number of functional evaluations required for MIMIC was also signifcnatly higher than the other algorithms, as was its training time. This was primarily due to the fact that MIMIC is signifciantly more complex to operate than the other algorithms. Overall, MIMIC was a very poor candidate for the traveling salesman problem.

Overall, the genetic algorithm proved to be the best strategy for the traveling salesman problem. It's ability to combine sub-journeys and paths into an optimal route was strongly advantageous. The results showed the strong benefit of mutations on this problem. Similarly, it's speed an relative simplicity were also very beneficial.

3) Flip flop

Overview

The flipflop problem is another common optimization, where one attempts to count the number of bits that alternate with its next neighbor in a bit string. This problem is particularly interesting because, since the strings are randomized, there is significant potential for a large number of local mimimum and maximums. It is also particularly well suited for algorithms that can handle pattern matching and previous knowledge well.
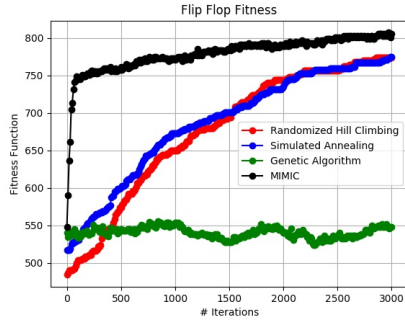
Fitness vs Iterations of 4 Optimization Strategies

Function Evals vs Iterations of 4 Optimization Strategies
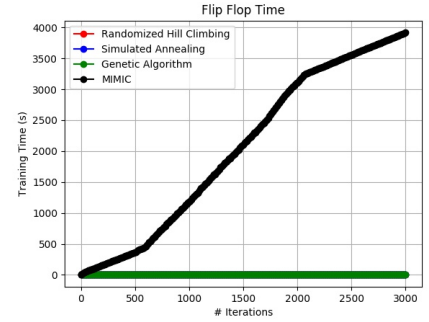
Training Time vs Iterations of 4 Optimization Strategies

Analysis

For this optimization problem, in order to ensure accuracy and a fair distribution, 5 different trials were used of 3000 iterations each. A bit string of length 1000 was used, where each bit was randomized as either 1 or 0. Each optimization strategy was allowed to evaluate over a parameter range in an attempt to find the optimal parameters for that strategy. The results are then summarized below.

For randomized hill climbing, there were no other parameters to tune. The algoritm was simply run for various iterations in a set of 5 trials. In terms of fitness function performance, it actually performed rather well–with it's performance increasing over time. Randomized hill climbing required very limited function evaluations and it trained quickly due to its simplicity. The reason for the stragegy's success can be attributed to the fact that random guessing is somewhat akin to the distribution of the input bit string in that the distribution fo the next neighbor should, on average, change.

Similar to the last problem, simulated annealing was very similar to randomized hill climbing. Cooling efficients of 0.15, 0.35, 0.55, 0.75, and 0.95 were used. It was found that the efficient 0.15 was the most effective in this case. Simulated annealing approach closely tracked the randomized hill climbing results yet again. Similar to hill climbing, the random variations of simulated annealing can roughly identify to the variations of the distribution of the input. Over time, as with randomized hill climbing, the performance of the strategy signficantly increased. Also Similar to randomized hill climbing, its' function evals count and training time were extremely low.

The genetic algorithm employed used a population size of 100 and various configurations of 50, 30, and 10 mutations/mate combinations. The graph above depicts the succes rates when using 30 mutation and 30 matings. The approach performed singificantly worse than the others. The genetic algorithm had a very tough time recognizing the pattern that the input tended to follow and the mutations and matings had very little effect on finding an optimal solution. While t number of evals was higher for the this strategy due to the amount of attempts at mutating and mating–it did not help to achieve a better result.

Iterations	Fitness	Training Time	Function Evals
0	548.0	10.292430561	1000
50	713.0	54.323489845	6500
100	745.0	89.781651075	12000
200	753.0	159.161412402	23000
300	755.0	227.402023603	34000
500	761.0	364.872433042	56000
750	771.0	741.847804258	83500
1000	771.0	1186.13889593	111000
1500	785.0	2107.83467502	166000
2000	789.0	3107.1784734	221000
2500	799.0	3557.31754274	276000
3000	805.0	3916.14983514	331000

MIMIC Results w/ .5 threshold

Mimic's parameters were set to 100 samples and a 'to-keep' number of 50. The threshold for the discrete dependency tree mimic used was sampled from the range 0.1, 0.3, 0.5, 0.7, and 0.9. MIMIC performed incredbly on this problem–mainly due to its ability to handle patterns very well and to keep track of historical knowledge. The number of evaluations required for MIMIC was definitely higher than the other algorithms, as was its training time, but its performance also far outweighed the other strategies. Overall, MIMIC proved incredibly succesful at the flip flop problem.

Conclusion

Above, four different machine learning optimization strategies were applied to 3 different optimization problems. Interestingly enough, each optimization problem had its own inherit properties that proved more suitable for solving for differnet optimization strategies.

The continous peaks problem, with its high amount of randomness and lack of predictable behaviour, proved to be optimized exceptionally well by simulated annealing. Simulated annealings ability to randomly sample nearby solutions with effective cooling thresholds were very helpful for this problem in that it avoided getting trapped in local optimums but reduced the change of skipping over more significant nearby optimums.

The traveling salesman problem, which, in my opinion, proved to be the most interesting, was most aptly solved by a genetic algorithm. Travelling salesman required various subjourneys to be optimized and combined to find an optimal path about various points. Whereas random and pattern solutions failed badly, the genetic algorithm was able to mutate and combine to form optimal solutions quite well.

The final problem, flip flop, had an input that was highly repetetive and tended to follow a patterned distribution. In such, optimization strategies that sought to solve based on subgroups of parameters or previous knowledge performed rather well, whereas strategies that were left to chance or random combinations performed poorly. MIMIC, with it's ability to solve patterns and repetition well, as well as its ability to retain prior experiences, proved to be far and away the most optimal solution for this problem.